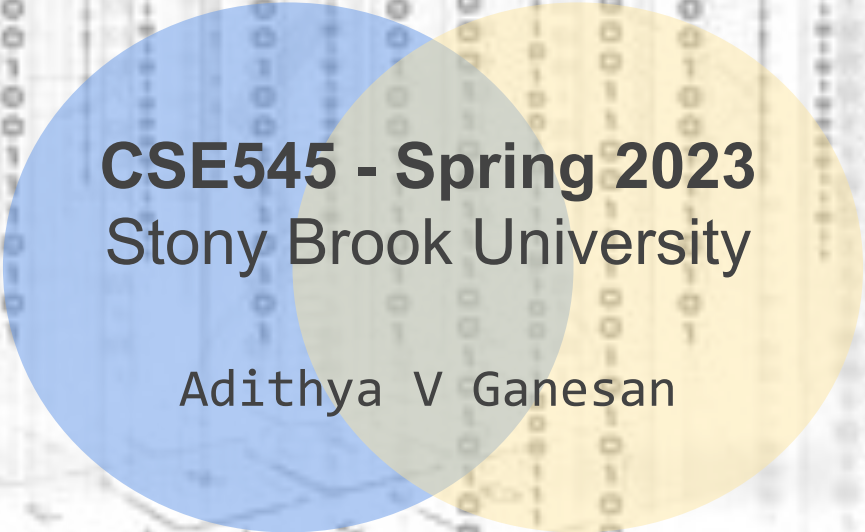


Neural Network Workflow Systems



CSE545 - Spring 2023
Stony Brook University

Adithya V Ganesan

with
PyTorch

Big Data Analytics, The Class

Goal: Generalizations
A model or summarization of the data.

Data Workflow Frameworks

Analytics and Algorithms

Hadoop File System✓
MapReduce✓
Streaming✓
Spark✓
Deep Learning Frameworks

Similarity Search
Hypothesis Testing
Transformers/Self-Supervision
Recommendation Systems
Link Analysis

Limitations of Spark

Spark is fast for being so flexible

- Fast: RDDs in memory + Lazy evaluation: optimized chain of operations.
- Flexible: Many transformations -- can contain any custom code.

Limitations of Spark

Spark is fast for being so flexible

- Fast: RDDs in memory + Lazy evaluation: optimized chain of operations.
- Flexible: Many transformations -- can contain any custom code.

However:

- Hadoop MapReduce can still be better for extreme IO, data that will not fit in memory across cluster.


Limitations of Spark

Spark is fast for being so flexible

- Fast: RDDs in memory + Lazy evaluation: optimized chain of operations.
- Flexible: Many transformations -- can contain any custom code.

However:

- Hadoop MapReduce can still be better for extreme IO, data that will not fit in memory across cluster.




Limitations of Spark

Spark is fast for being so flexible

- Fast: RDDs in memory + Lazy evaluation: optimized chain of operations.
- Flexible: Many transformations -- can contain any custom code.

However:

- Hadoop MapReduce can still be better for extreme IO, data that will not fit in memory across cluster.
- Modern machine learning (esp. Deep learning), a common big data task, requires heavy numeric computation.



* this is the subjective approximation of the instructor as of February 2020. A lot of factors at play.


Limitations of Spark

Spark is fast for being so flexible

- Fast: RDDs in memory + Lazy evaluation: optimized chain of operations.
- Flexible: Many transformations -- can contain any custom code.

However:

- Hadoop MapReduce can still be better for extreme IO, data that will not fit in memory across cluster.
- Modern machine learning (esp. Deep learning), a common big data task, requires heavy numeric computation.



* this is the subjective approximation of the instructor as of February 2020. A lot of factors at play.

Learning Objectives

- Understand a neural network as transformations on tensors.
- Understand PyTorch as a data workflow system.
 - Know the key components of PyTorch
 - Understand the key concepts around *distributed* neural network processing in PyTorch.
- Establish a foundation to distribute deep learning models

Linear Regression

Linear Regression: $\hat{y} = \beta X$

Objective: *Learn w , such that $(y - \beta X)^2$ is minimized*

Linear Regression

Linear Regression: $\hat{y} = \beta X$

Objective: *Learn w , such that $(y - \beta X)^2$ is minimized*

How do we solve for β ?

Linear Regression

Linear Regression: $\hat{y} = \beta X$

Objective: *Learn w , such that $(y - \beta X)^2$ is minimized*

How do we solve for β ?

1. Analytic Gradient: Differentiate the objective, solve the system of equations by equating it to 0

Linear Regression

Linear Regression: $\hat{y} = \beta X$

Objective: *Learn w , such that $(y - \beta X)^2$ is minimized*

How do we solve for β ?

1. Analytic Gradient: Differentiate the objective, solve the system of equations by equating it to 0

$$\beta_{opt} = (X^T X)^{-1} X^T y$$

Linear Regression

Linear Regression: $\hat{y} = \beta X$

Objective: *Learn w , such that $(y - \beta X)^2$ is minimized*

How do we solve for β ?

1. Analytic Gradient: Differentiate the objective, solve the system of equations by equating it to 0
2. Numerical Gradient: Start at a random point and move in the direction of minima until optima is reached


Linear Regression

Linear Regression: $\hat{y} = \beta X$

Objective: *Learn w , such that $(y - \beta X)^2$ is minimized*

How do we solve for β ?

1. Analytic Gradient: Differentiate the objective, solve the system of equations by equating it to 0
2. **Numerical Gradient: Start at a random point and move in the direction of minima until optima is reached**



Numerical Gradient Approach

Linear Regression: Trying to find “betas” that minimize:

$$\hat{\beta} = \operatorname{argmin}_{\beta} \left\{ \sum_i^N (y_i - \hat{y}_i)^2 \right\}$$

Numerical Gradient Approach

Linear Regression: Trying to find “betas” that minimize:

$$\hat{\beta} = \operatorname{argmin}_{\beta} \left\{ \sum_i^N (y_i - \hat{y}_i)^2 \right\}$$

matrix multiply

$$\hat{y}_i = X_i \beta$$

Numerical Gradient Approach

Linear Regression: Trying to find “betas” that minimize:

$$\hat{\beta} = \operatorname{argmin}_{\beta} \left\{ \sum_i^N (y_i - \hat{y}_i)^2 \right\}$$

matrix multiply

$$\hat{y}_i = X_i \beta$$

Thus:

$$\hat{\beta} = \operatorname{argmin}_{\beta} \left\{ \sum_{i=0}^N (y_i - X_i \beta)^2 \right\}$$

Numerical Gradient Approach

Linear Regression: Trying to find “betas” that minimize:

$$\hat{\beta} = \operatorname{argmin}_{\beta} \left\{ \sum_i^N (y_i - \hat{y}_i)^2 \right\}$$

$$\hat{y}_i = X_i \beta \quad \text{Thus:} \quad \hat{\beta} = \operatorname{argmin}_{\beta} \left\{ \sum_{i=0}^N (y_i - X_i \beta)^2 \right\}$$

How to update? $\beta_{new} = \beta_{prev} - \alpha * \text{grad}$

Numerical Gradient Approach

Linear Regression: Trying to find “betas” that minimize:

$$\hat{\beta} = \operatorname{argmin}_{\beta} \left\{ \sum_i^N (y_i - \hat{y}_i)^2 \right\}$$

$$\hat{y}_i = X_i \beta \quad \text{Thus:} \quad \hat{\beta} = \operatorname{argmin}_{\beta} \left\{ \sum_{i=0}^N (y_i - X_i \beta)^2 \right\}$$

How to update? $\beta_{new} = \beta_{prev} - \alpha * \text{grad}$



α : Learning Rate

Numerical Gradient Approach

Linear Regression: Trying to find “betas” that minimize:

$$\hat{\beta} = \operatorname{argmin}_{\beta} \left\{ \sum_i^N (y_i - \hat{y}_i)^2 \right\}$$


$$\hat{y}_i = X_i \beta$$

Thus:

$$\hat{\beta} = \operatorname{argmin}_{\beta} \left\{ \sum_{i=0}^N (y_i - \right.$$

How to update? $\beta_{new} = \beta_{prev} - \alpha * \text{grad}$

α : Learning Rate



Numerical Gradient Approach

Linear Regression: Trying to find “betas” that minimize:

Gradient Descent: $\beta_{new} = \beta_{prev} - \alpha * \text{grad}$

Numerical Gradient Approach

Linear Regression: Trying to find “betas” that minimize:

Gradient Descent: $\beta_{new} = \beta_{prev} - \alpha * \text{grad}$


But there are other gradient descent based optimization methods which are better*

Numerical Gradient Approach

Linear Regression: Trying to find “betas” that minimize:

Gradient Descent: $\beta_{new} = \beta_{prev} - \alpha * \text{grad}$

But there are other gradient descent based op



Animation: Alec Radford

Linear Regression as DAG


How do Machine learning/ Deep learning frameworks represent these models?

Linear Regression as DAG

How do Machine learning/ Deep learning frameworks represent these models?

Computational Graph!


Linear Regression as DAG




$$L = (y - \beta x)^2$$

Activations


Rectified linear unit (ReLU): $ReLU(z) = \max(0, z)$




Hyperbolic tangent: $\tanh(z) = (e^{2z} - 1) / (e^{2z} + 1)$



Linear Regression as DAG



Linear Regression as DAG



```
import torch
from torch import nn

x = torch.Tensor(input)
beta = torch.random.randn(X.shape, 1)
z = torch.matmul(x, beta)
yhat = nn.functional.relu(z)
loss = nn.MSELoss(yhat, torch.Tensor(y))
```

PyTorch Demo

Native Linear Regression Implementation ([Link](#))

Torch.nn Linear Regression Implementation ([Link](#))

How to train GPT3?

Time to train Bert Large (330 M) on K80, which is 530 times smaller than GPT3

# GPUs	Training Time (minutes)	Per-GPU Scaling Efficiency
1	399	1.00

How to train GPT3?

Time to train Bert Large (330 M) on K80, which is 530 times smaller than GPT3

# GPUs	Training Time (minutes)	Per-GPU Scaling Efficiency
1	399	1.00

For the same amount of data, GPT3 can be trained in 212k mins = 3533 hours = 147 days*

**GPT3 wont fit into the memory of a single K80*

How to train GPT3?

Time to train Bert Large (330 M) on K80, which is 530 times smaller than GPT3


# GPUs	Training Time (minutes)	Per-GPU Scaling Efficiency
1	399	1.00
2	214	0.93
4	118	0.85
8	61	0.82

Distributed Training

- Parallelism :
 - Data Parallelism
 - Model Parallelism
 - Hybrid

Distributed PyTorch Training


- Data Parallelism: Scatter dataset into parts across different workers to train on subsets and sync gradients



Data Parallelism


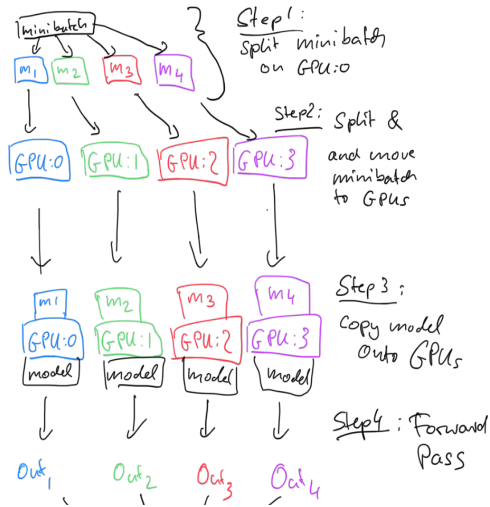
Distributed PyTorch Training

- Data Parallelism: Scatter dataset into parts across different workers to train on subsets and sync gradients
- Modes of Data Parallelism :
 - DataParallel
 - DistributedDataParallel



Distributed PyTorch Training

Data Parallel: How it works?

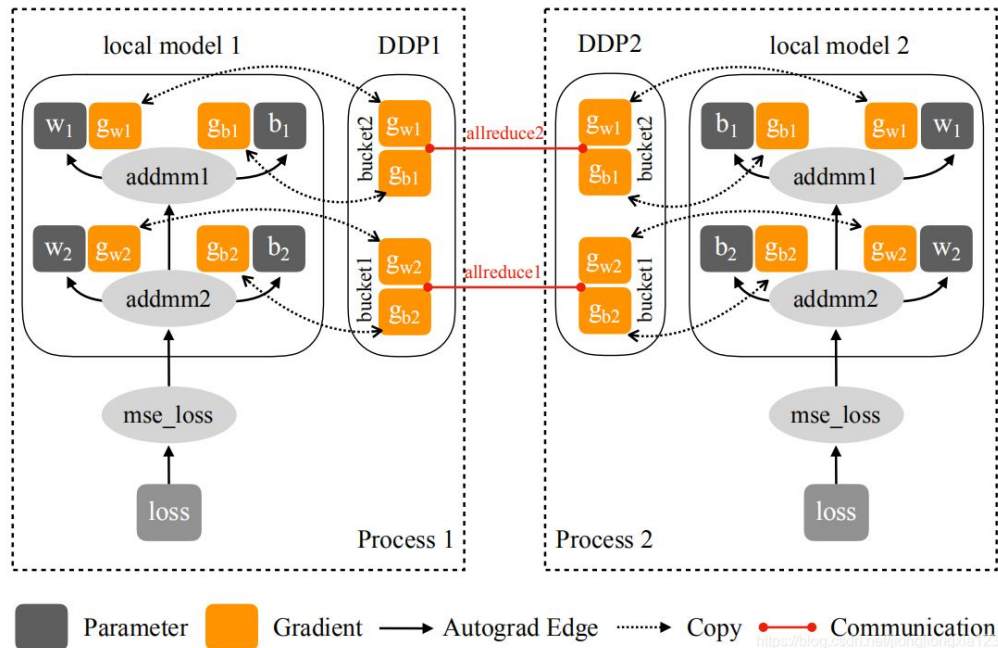


Distributed PyTorch Training

- Data Parallel
 - Most simple form of parallelism with minimal code change
 - Downside: Slower form of parallelism - involves inter node communication 3x per training step

Distributed PyTorch Training

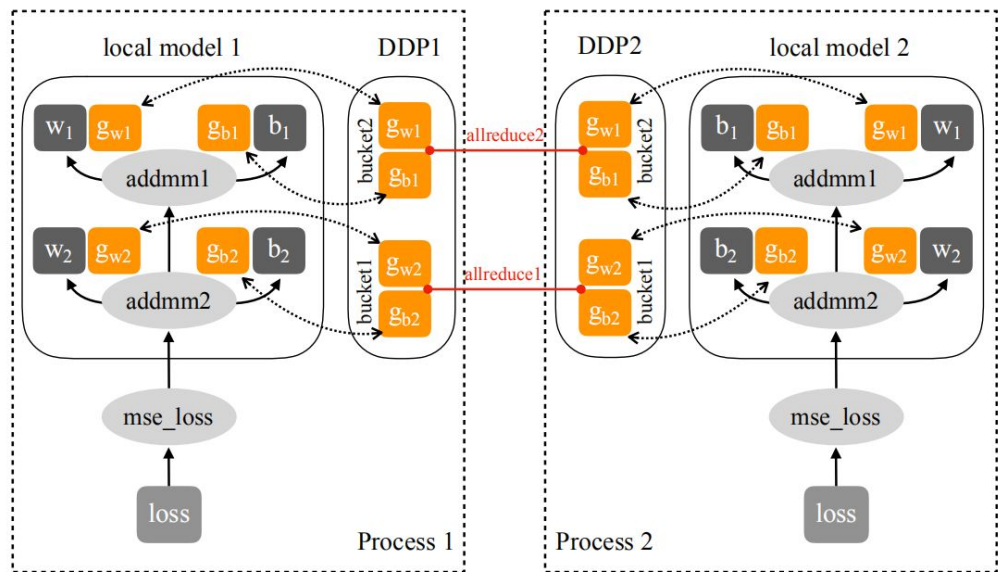
DistributedDataParallel: How it works?




([Li et al., 2020](#))

Distributed PyTorch Training

- DistributedDataParallel



■ Parameter ■ Gradient → Autograd Edge Copy —●— Communication



AllReduce

([Li et al., 2020](#))

Distributed PyTorch Training

- DistributedDataParallel ([Li et al., 2020](#))
 - Efficient form of parallelism but involves a little extra code change*
 - Performs AllReduce on the computed gradients across all nodes and machines

** Extra code change if you are implementing using Pytorch. It has been made extremely simple by*

Distributed PyTorch Training

- DistributedDataParallel ([Li et al., 2020](#))
 - Efficient form of parallelism but involves a little extra code change*
 - Performs AllReduce on the computed gradients across all nodes and machines
 - Downside: Python pickles all objects while spawning multiple processes (which happens in DDP). Code might crash if an object is not pickle-able

** Extra code change if you are implementing using Pytorch. It has been made extremely simple by*

Distributed PyTorch Training


Model Parallelism: Distribute layer(s) of the model into different machines/GPUs to train a very large network.

Distributed PyTorch Training

- Model Parallelism: Distribute layer(s) of the model into different machines/GPUs to train a very large network.
- Model Parallelism
 - Naive Model Parallelism
 - Pipelined Parallelism

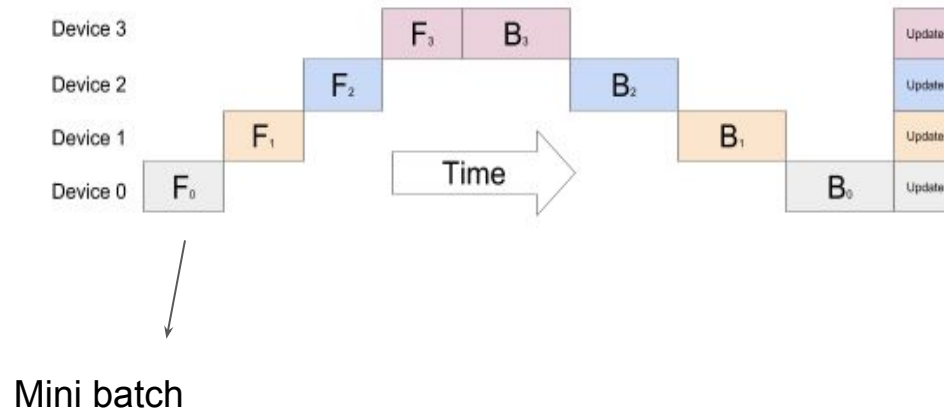
Distributed PyTorch Training

- Naive Model Parallelism



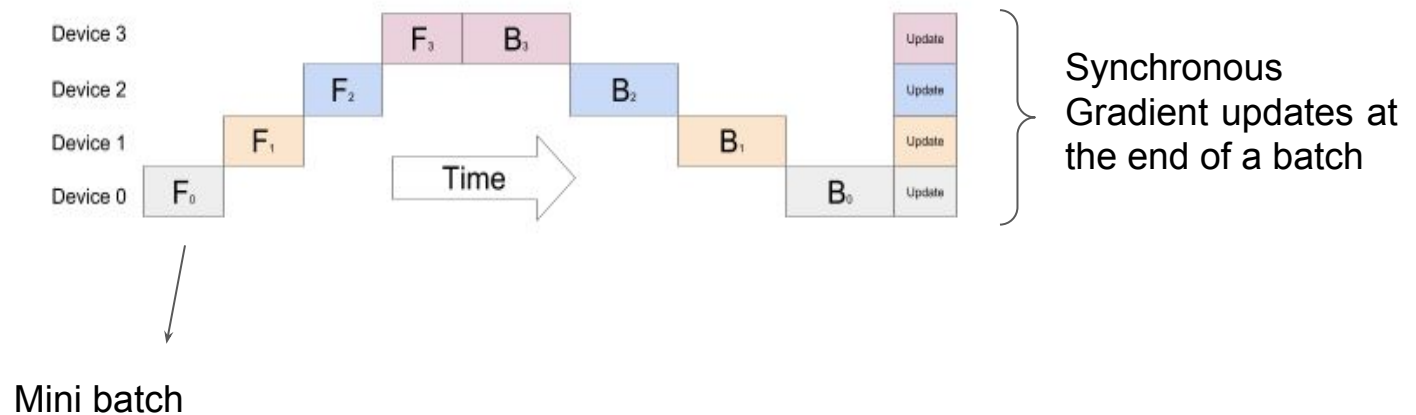
Distributed PyTorch Training

- Naive Model Parallelism



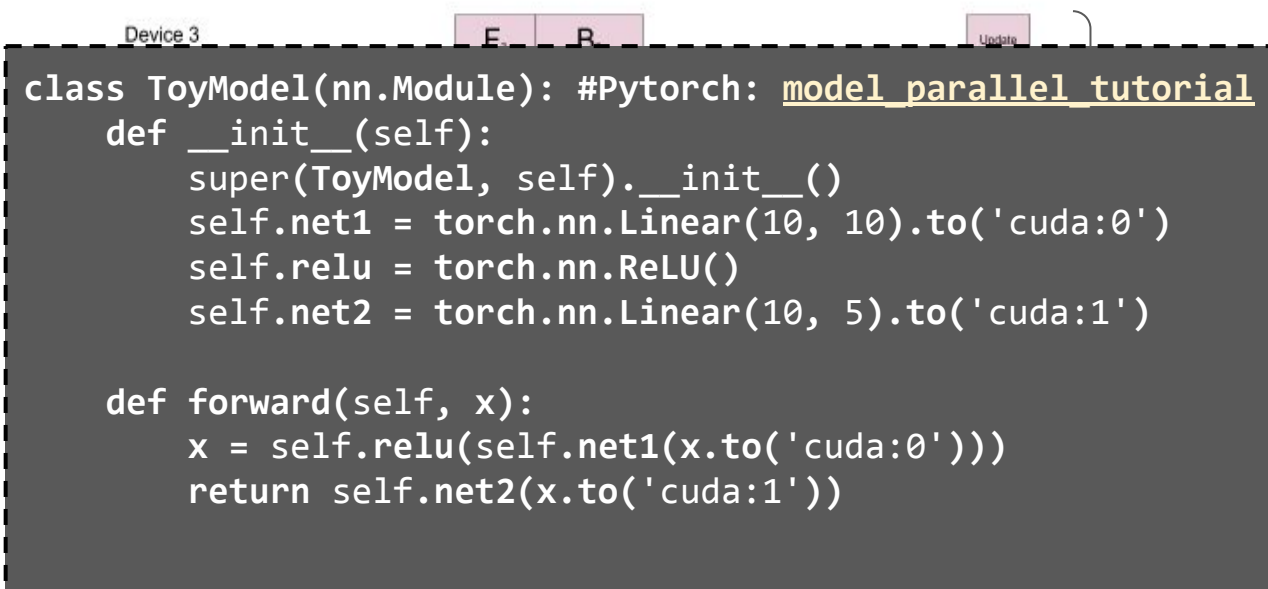
Distributed PyTorch Training

- Naive Model Parallelism



Distributed PyTorch Training

- Naive Model Parallelism



The diagram illustrates Naive Model Parallelism. It features a large dark gray box with a dashed border containing Python code for a `ToyModel` class. Above the box, there are labels for 'Device 3', 'E', 'B', and 'Update'. To the right of the box, there is a label 'Continuous updates at the end of a batch'.

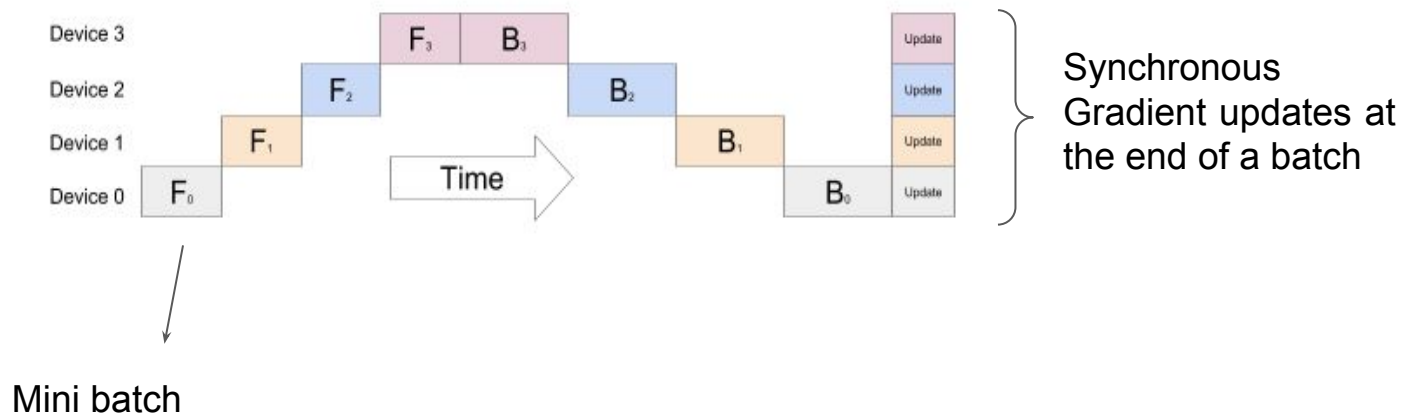
```
class ToyModel(nn.Module): #Pytorch: model_parallel_tutorial
    def __init__(self):
        super(ToyModel, self).__init__()
        self.net1 = torch.nn.Linear(10, 10).to('cuda:0')
        self.relu = torch.nn.ReLU()
        self.net2 = torch.nn.Linear(10, 5).to('cuda:1')

    def forward(self, x):
        x = self.relu(self.net1(x.to('cuda:0')))
        return self.net2(x.to('cuda:1'))
```

Continuous
updates at
the end of a batch

Distributed PyTorch Training


- Naive Model Parallelism



Severe under utilization of resources due to sequential dependency of the network


Distributed PyTorch Training

- Pipelined Parallelism



Distributed PyTorch Training

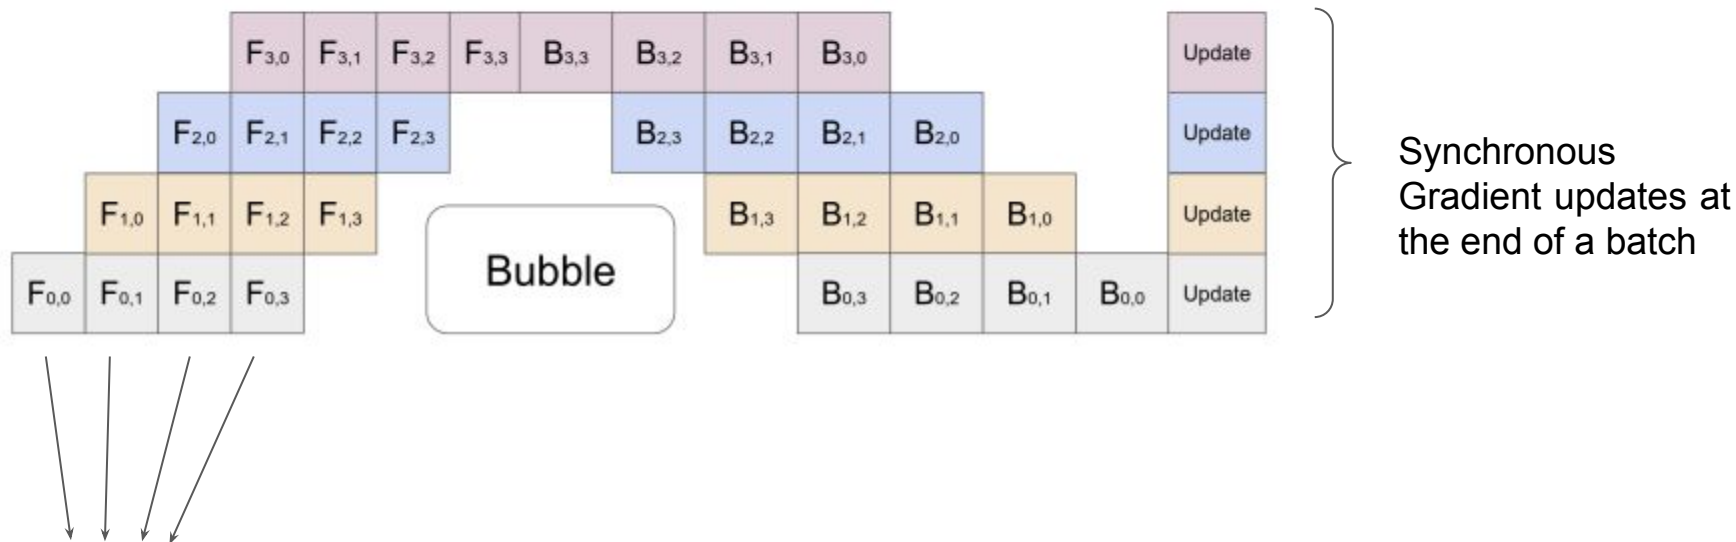
- Pipelined Parallelism



Mini batch split into
micro batches

Distributed PyTorch Training


- Pipelined Parallelism



Mini batch split into
micro batches

Distributed PyTorch Training

- Pipelined Parallelism




Mini batch split into
micro batches

Provides high utilization of workers while ensuring reliable + stable training

Distributed PyTorch Training

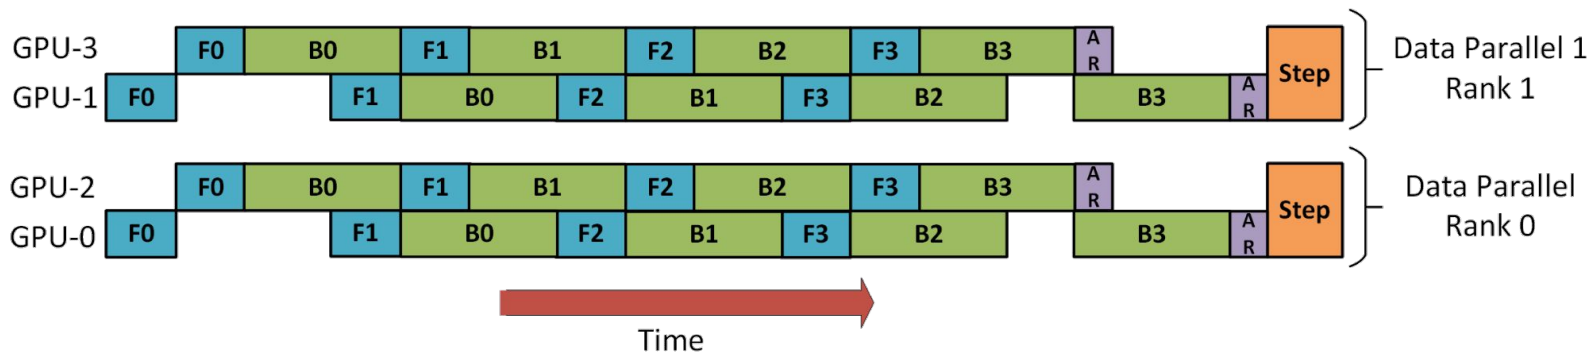
- Pipelined Parallelism



PyTorch: [Model Parallel best practices](#)

Distributed PyTorch Training

- Hybrid
 - DeepSpeed ([Rasley et al., 2020](#))



Horovod: PyTorch 🤝 PySpark

Horovod is a distributed deep learning training framework.

Horovod helps scaling single GPU (worker) into multi-GPU or even multi-host training without no code change

Horovod on [spark](#): “provides a convenient wrapper around Horovod that makes running distributed training jobs in Spark clusters easy”

